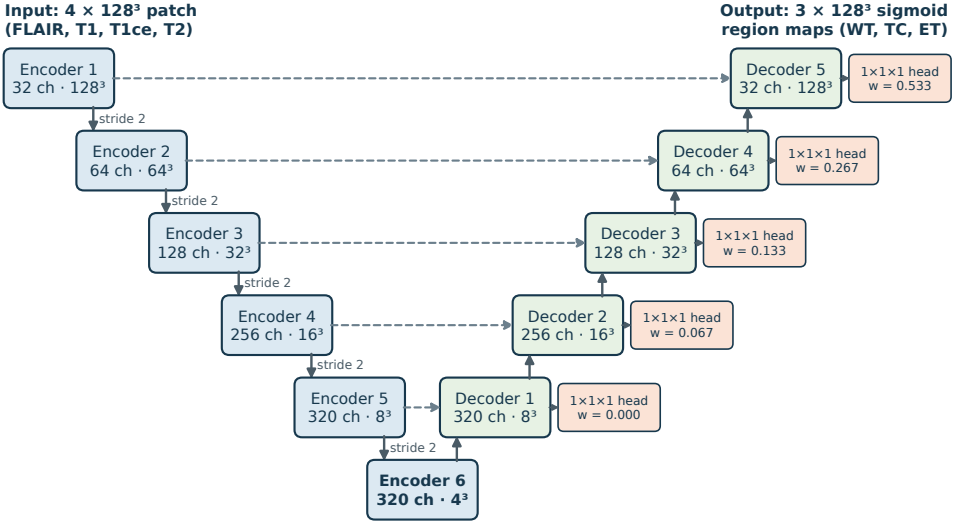

Each stage: 2 × [3×3×3 Conv3d → InstanceNorm3d → LeakyReLU]; stride 1 in stage 1, stride 2 (first convolution) in stages 2–6.

Decoder: transposed convolution (kernel = stride = 2), concatenation with the skip connection (dashed), then 2 convolutions.

Orange: deep-supervision heads used in training only; loss weights $1/2^i$ normalised to sum to 1, lowest-resolution weight set to 0.